

Economics 103 – Statistics for Economists

Francis J. DiTraglia

University of Pennsylvania

Lecture # 9

Discrete RVs – Part II

Variance and Standard Deviation of a RV

The Defs are the same for continuous RVs, but the method of calculating will differ.

Variance (Var)

$$\sigma^2 = \text{Var}(X) = E[(X - \mu)^2] = E[(X - E[X])^2]$$

Standard Deviation (SD)

$$\sigma = \sqrt{\sigma^2} = \text{SD}(X)$$

Key Point

Variance and std. dev. are *expectations of functions of a RV*

It follows that:

1. Variance and SD are constants
2. To derive facts about them you can use the facts you know about expected value

How To Calculate Variance for Discrete RV?

Remember: it's just a function of X !

$$\text{Recall that } \mu = E[X] = \sum_{\text{all } x} xp(x)$$

$$\text{Var}(X) = E[(X - \mu)^2] = \sum_{\text{all } x} (x - \mu)^2 p(x)$$

Shortcut Formula For Variance

This is *not* the definition, it's a shortcut for doing calculations:

$$\text{Var}(X) = E[(X - \mu)^2] = E[X^2] - (E[X])^2$$

We'll prove this in an upcoming lecture.

Example: The Shortcut Formula



Let $X \sim \text{Bernoulli}(1/2)$.

1. What is $E[X]$? $E[X] = 0 \times 1/2 + 1 \times 1/2 = 1/2$
2. What is $E[X^2]$? $E[X^2] = 0^2 \times 1/2 + 1^2 \times 1/2 = 1/2$
3. What is $\text{Var}(X)$? $E[X^2] - (E[X])^2 = 1/2 - (1/2)^2 = 1/4$

Variance of Bernoulli RV – via the Shortcut Formula

Step 1 – $E[X]$

$$\mu = E[X] = \sum_{x \in \{0,1\}} p(x) \cdot x = (1-p) \cdot 0 + p \cdot 1 = p$$

Step 2 – $E[X^2]$

$$E[X^2] = \sum_{x \in \{0,1\}} x^2 p(x) = 0^2(1-p) + 1^2 p = p$$

Step 3 – Combine with Shortcut Formula

$$\sigma^2 = \text{Var}[X] = E[X^2] - (E[X])^2 = p - p^2 = p(1-p)$$

Variance of Bernoulli RV – Without Shortcut

You will fill in the missing steps on Problem Set 5.

$$\begin{aligned}\sigma^2 &= \text{Var}(X) = \sum_{x \in \{0,1\}} (x - \mu)^2 p(x) \\ &= \sum_{x \in \{0,1\}} (x - p)^2 p(x) \\ &\vdots \\ &= p(1 - p)\end{aligned}$$

Variance of a Linear Function



Suppose X is a random variable with $\text{Var}(X) = \sigma^2$ and a, b are constants. What is $\text{Var}(a + bX)$?

- (a) σ^2
- (b) $a + \sigma^2$
- (c) $b\sigma^2$
- (d) $a + b\sigma^2$
- (e) $b^2\sigma^2$

Variance and SD are *NOT* Linear

$$\text{Var}(a + bX) = b^2\sigma^2$$

$$\text{SD}(a + bX) = |b|\sigma$$

These should look familiar from the related results for sample variance and std. dev. that you worked out on an earlier problem set.

Variance of a Linear Transformation

$$\begin{aligned}\text{Var}(a + bX) &= E \left[\{(a + bX) - E(a + bX)\}^2 \right] \\&= E \left[\{(a + bX) - (a + bE[X])\}^2 \right] \\&= E \left[(bX - bE[X])^2 \right] \\&= E[b^2(X - E[X])^2] \\&= b^2 E[(X - E[X])^2] \\&= b^2 \text{Var}(X) = b^2 \sigma^2\end{aligned}$$

The key point here is that variance is defined in terms of expectation and expectation is linear.

Binomial Random Variable

What we get if we sum a bunch of indep. Bernoulli RVs

Binomial Random Variable

Let X = the sum of n independent Bernoulli trials, each with probability of success p . Then we say that: $X \sim \text{Binomial}(n, p)$

Parameters

p = probability of “success,” n = # of trials

Support


$\{0, 1, 2, \dots, n\}$

Probability Mass Function (pmf)

$$p(x) = \binom{n}{x} p^x (1 - p)^{n-x}$$

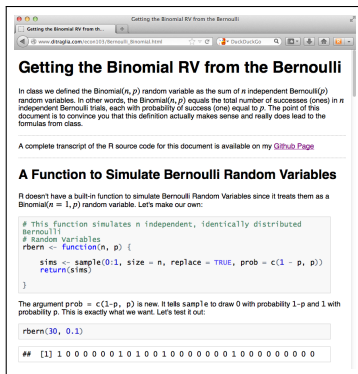
http://fditraglia.shinyapps.io/binom_cdf/

Try playing around with all three sliders. If you set the second to 1 you get a Bernoulli.



http://fditraglia.github.com/Econ103Public/Rtutorials/Bernoulli_Binomial.html

Source Code on my [Github Page](#)



Getting the Binomial RV from the Bernoulli

In class we defined the $\text{Binomial}(n, p)$ random variable as the sum of n independent $\text{Bernoulli}(p)$ random variables. In other words, the $\text{Binomial}(n, p)$ equals the total number of successes (ones) in n independent Bernoulli trials, each with probability of success (one) equal to p . The point of this document is to convince you that this definition actually makes sense and really does lead to the formulas from class.

A complete transcript of the R source code for this document is available on my [Github Page](#).

A Function to Simulate Bernoulli Random Variables

R doesn't have a built-in function to simulate Bernoulli Random Variables since it treats them as a $\text{Binomial}(n = 1, p)$ random variable. Let's make our own:

```
# This function simulates n independent, identically distributed
# Bernoulli
# Random Variables
rbern <- function(n, p) {
  sims <- sample(0:1, size = n, replace = TRUE, prob = c(1 - p, p))
  return(sims)
}
```

The argument `prob = c(1-p, p)` is new. It tells `sample` to draw 0 with probability $1-p$ and 1 with probability p . This is exactly what we want. Let's test it out:

```
rbern(30, 0.1)
```

```
## [1] 1 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0
```

Don't forget this!

Binomial RV counts up the *total* number of successes (ones) in n indep. Bernoulli trials, each with prob. of success p .

Where does the Binomial pmf come from?

Question

Suppose we flip a fair coin 3 times. What is the probability that we get exactly 2 heads?

Answer

Three basic outcomes make up this event: $\{HHT, HTH, THH\}$, each has probability $1/8 = 1/2 \times 1/2 \times 1/2$. Basic outcomes are mutually exclusive, so sum to get $3/8 = 0.375$

Where does the Binomial pmf come from?

Question

Suppose we flip an *unfair* coin 3 times, where the probability of heads is $1/3$. What is the probability that we get exactly 2 heads?

Answer

No longer true that *all* basic outcomes are equally likely, but those with exactly two heads *still are*

$$P(HHT) = (1/3)^2(1 - 1/3) = 2/27$$

$$P(THH) = 2/27$$

$$P(HTH) = 2/27$$

Summing gives $2/9 \approx 0.22$

Where does the Binomial pmf come from?

Starting to see a pattern?

Suppose we flip an unfair coin 4 times, where the probability of heads is $1/3$. What is the probability that we get exactly 2 heads?

HHTT TTHH

HTHT THTH

HTTH THTT

Six equally likely, mutually exclusive
basic outcomes make up this event:

$$\binom{4}{2} (1/3)^2 (2/3)^2$$

Multiple RVs *at once* - Definition of Joint PMF


Let X and Y be discrete random variables. The joint probability mass function $p_{XY}(x, y)$ gives the probability of each pair of realizations (x, y) in the support:

$$p_{XY}(x, y) = P(X = x \cap Y = y)$$

Example: Joint PMF in Tabular Form

		Y		
		1	2	3
X	0	1/8	0	0
	1	0	1/4	1/8
	2	0	1/4	1/8
	3	1/8	0	0

Plot of Joint PMF



What is $p_{XY}(1, 2)$?



		Y		
		1	2	3
X	0	1/8	0	0
	1	0	1/4	1/8
	2	0	1/4	1/8
	3	1/8	0	0

$$p_{XY}(1, 2) = P(X = 1 \cap Y = 2) = 1/4$$

What is $p_{XY}(2, 1)$?



		Y		
		1	2	3
X	0	1/8	0	0
	1	0	1/4	1/8
	2	0	1/4	1/8
	3	1/8	0	0

$$p_{XY}(2, 1) = P(X = 2 \cap Y = 1) = 0$$

Properties of Joint PMF

1. $0 \leq p_{XY}(x, y) \leq 1$ for any pair (x, y)
2. The sum of $p_{XY}(x, y)$ over all pairs (x, y) in the support is 1:

$$\sum_x \sum_y p(x, y) = 1$$

Does this satisfy the properties of a joint pmf?



(A = YES, B = NO)

		Y		
		1	2	3
X	0	1/8	0	0
	1	0	1/4	1/8
	2	0	1/4	1/8
	3	1/8	0	0

1. $p(x, y) \geq 0$ for all pairs (x, y)
2. $\sum_x \sum_y p(x, y) = 1/8 + 1/4 + 1/8 + 1/4 + 1/8 + 1/8 = 1$

Joint versus Marginal PMFs

Joint PMF

$$p_{XY}(x, y) = P(X = x \cap Y = y)$$

Marginal PMFs

$$p_X(x) = P(X = x)$$

$$p_Y(y) = P(Y = y)$$

You can't calculate a joint pmf from marginals alone but you *can* calculate marginals from the joint!

Marginals from Joint

$$p_X(x) = \sum_{\text{all } y} p_{XY}(x, y)$$

$$p_Y(y) = \sum_{\text{all } x} p_{XY}(x, y)$$

Why?

$$\begin{aligned} p_Y(y) &= P(Y = y) = P\left(\bigcup_{\text{all } x} \{X = x \cap Y = y\}\right) \\ &= \sum_{\text{all } x} P(X = x \cap Y = y) = \sum_{\text{all } x} p_{XY}(x, y) \end{aligned}$$

To get the marginals sum “into the margins” of the table.

		Y			
		1	2	3	
X	0	1/8	0	0	1/8
	1	0	1/4	1/8	3/8
	2	0	1/4	1/8	3/8
	3	1/8	0	0	1/8
					1

$$p_X(0) = 1/8 + 0 + 0 = 1/8$$

$$p_X(1) = 0 + 1/4 + 1/8 = 3/8$$

$$p_X(2) = 0 + 1/4 + 1/8 = 3/8$$

$$p_X(3) = 1/8 + 0 + 0 = 1/8$$

What is $p_Y(2)$?



		Y			
		1	2	3	
X	0	1/8	0	0	
	1	0	1/4	1/8	
	2	0	1/4	1/8	
	3	1/8	0	0	
		1/4	1/2	1/4	1

$$p_Y(1) = 1/8 + 0 + 0 + 1/8 = 1/4$$

$$p_Y(2) = 0 + 1/4 + 1/4 + 0 = 1/2$$

$$p_Y(3) = 0 + 1/8 + 1/8 + 0 = 1/4$$